

PRACE GPU Programming Course

Alessio Sclocco and Ben van Werkhoven

December 16, 2022

netherlands
eSciencecenter

Introduction: Alessio Sclocco

Alessio Sclocco

eScience Research Engineer

Netherlands eScience Center

a.sclocco@esciencecenter.nl



Background:

- 2011-2012 junior researcher at VU Amsterdam
 - Working on GPUs for radio astronomy
- 2012-2017 PhD "Accelerating Radio Astronomy with Auto-Tuning" at VU Amsterdam
 - Under the supervision of professors Henri Bal and Rob van Nieuwpoort
- 2015-2016 scientific programmer at ASTRON, the Netherlands Institute for Radio Astronomy
 - Designing and developing a real-time GPU pipeline for the Westerbork radio telescope
- 2019 visiting scholar at Nanyang Technological University in Singapore
- 2017-2022 eScience Research Engineer at the Netherlands eScience Center
 - Radio astronomy, climate modeling, biology, natural language processing, high-energy physics

Introduction: Ben van Werkhoven

Ben van Werkhoven
Senior Research Engineer
Netherlands eScience Center
b.vanwerkhoven@esciencecenter.nl



Background:

- 2010-2014 PhD “Scientific Supercomputing with Graphics Processing Units” at the VU University Amsterdam in the group of prof. Henri Bal
- 2014-now working at the Netherlands eScience Center as the GPU expert in many different scientific research projects

GPU Programming since early 2009, worked on applications in computer vision, digital forensics, climate modeling, particle physics, geospatial databases, radio astronomy, and localization microscopy

Schedule

- 14:15 – 14:25 Course introduction
- 14:25 – 14:40 Introduction to GPU programming
- 14:40 – 14:50 GPU-Enabled libraries
- 14:50 – 15:05 Introduction to CUDA programming
- 15:05 – 15:20 Hands-on exercise

- 15:20 – 15:35 Break

- 15:35 – 15:50 CUDA memories part 1
- 15:50 – 16:05 Hands-on exercise
- 16:05 – 16:20 CUDA memories part 2
- 16:20 – 16:50 Hands-on exercise

- 16:50 - 17:05 Break

- 17:05 – 17:45 CUDA Program execution
- 17:45 – 17:55 Closing


- Get your own copy of the slides so you can read along and click on links
 - See: <https://github.com/benvanwerkhoven/gpu-course/>
 - Clone the repository to have access to slides and hands-on exercises
- Our slides are sometimes very wordy, this is intentional, so they may serve as a reference that you can read again later
- In code samples on the slides we sometimes abbreviate the code a bit to save space


Introduction to GPU Computing



What is a GPU?

- Graphics Processing Unit: the computing chip on a graphics card
- GPGPU: General Purpose computing on GPUs










Performance Comparison: CPUs vs GPUs





The World's Fastest Supercomputers: Oak Ridge's Frontier

- Number 1 in TOP500 list (Nov 2022)
 - 1.6 ExaFLOP/s peak
 - 1,000 cabinets
 - 21 MW power consumption
 - CPUs
 - AMD EPYC 64C 2GHz
 - GPUs
 - AMD Instinct MI250X
 - 8,730,112 cores
- Seven systems in top 10 using GPUs



- Different goals produce different designs
 - GPU assumes workload is highly parallel
 - CPU must be good at everything, parallel or not
- CPU: minimize latency experienced by 1 thread
 - Big on-chip caches
 - Sophisticated control logic
- GPU: maximize throughput of all threads
 - Multithreading can hide latency, so no big caches
 - Control logic
 - Much simpler
 - Less: share control logic across many threads




Why is GPU Programming different?

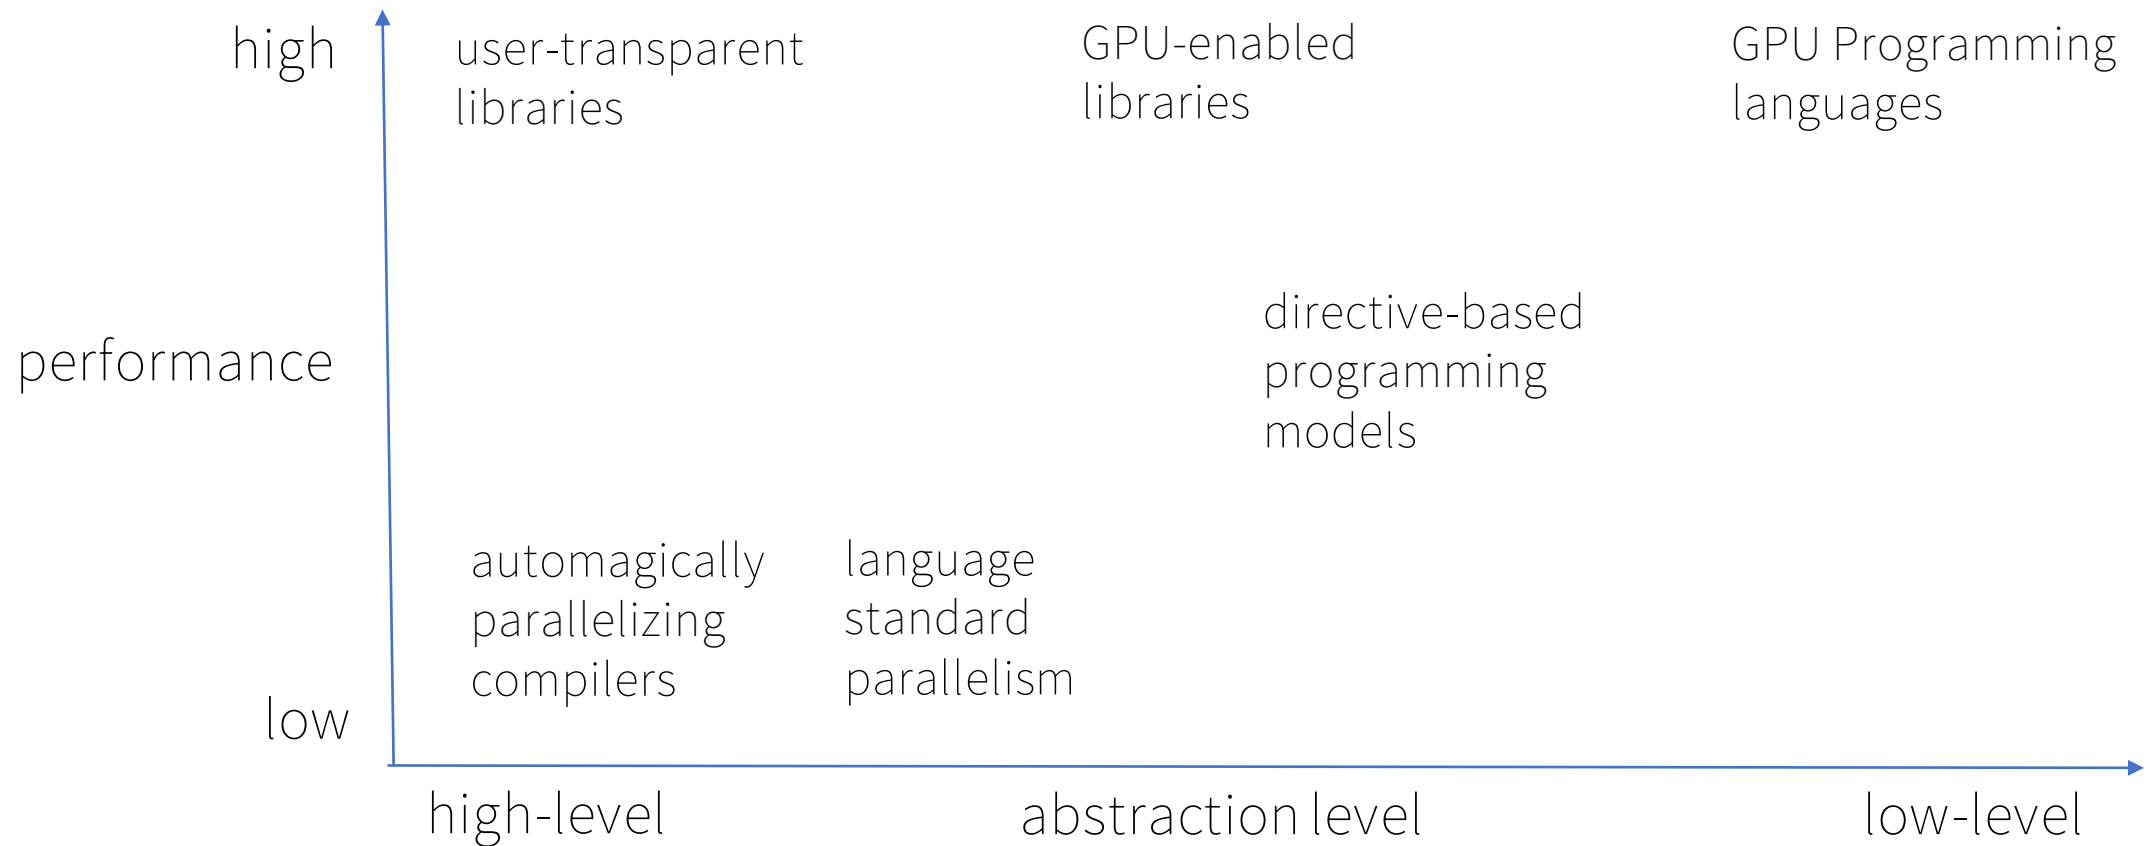
The computer architecture is very different:

- Algorithms need to be parallelized and mapped to the hardware
- Requires software to be rewritten in specialized programming language
- Optimizing for compute performance requires knowledge about hardware

GPUs are on separate devices:

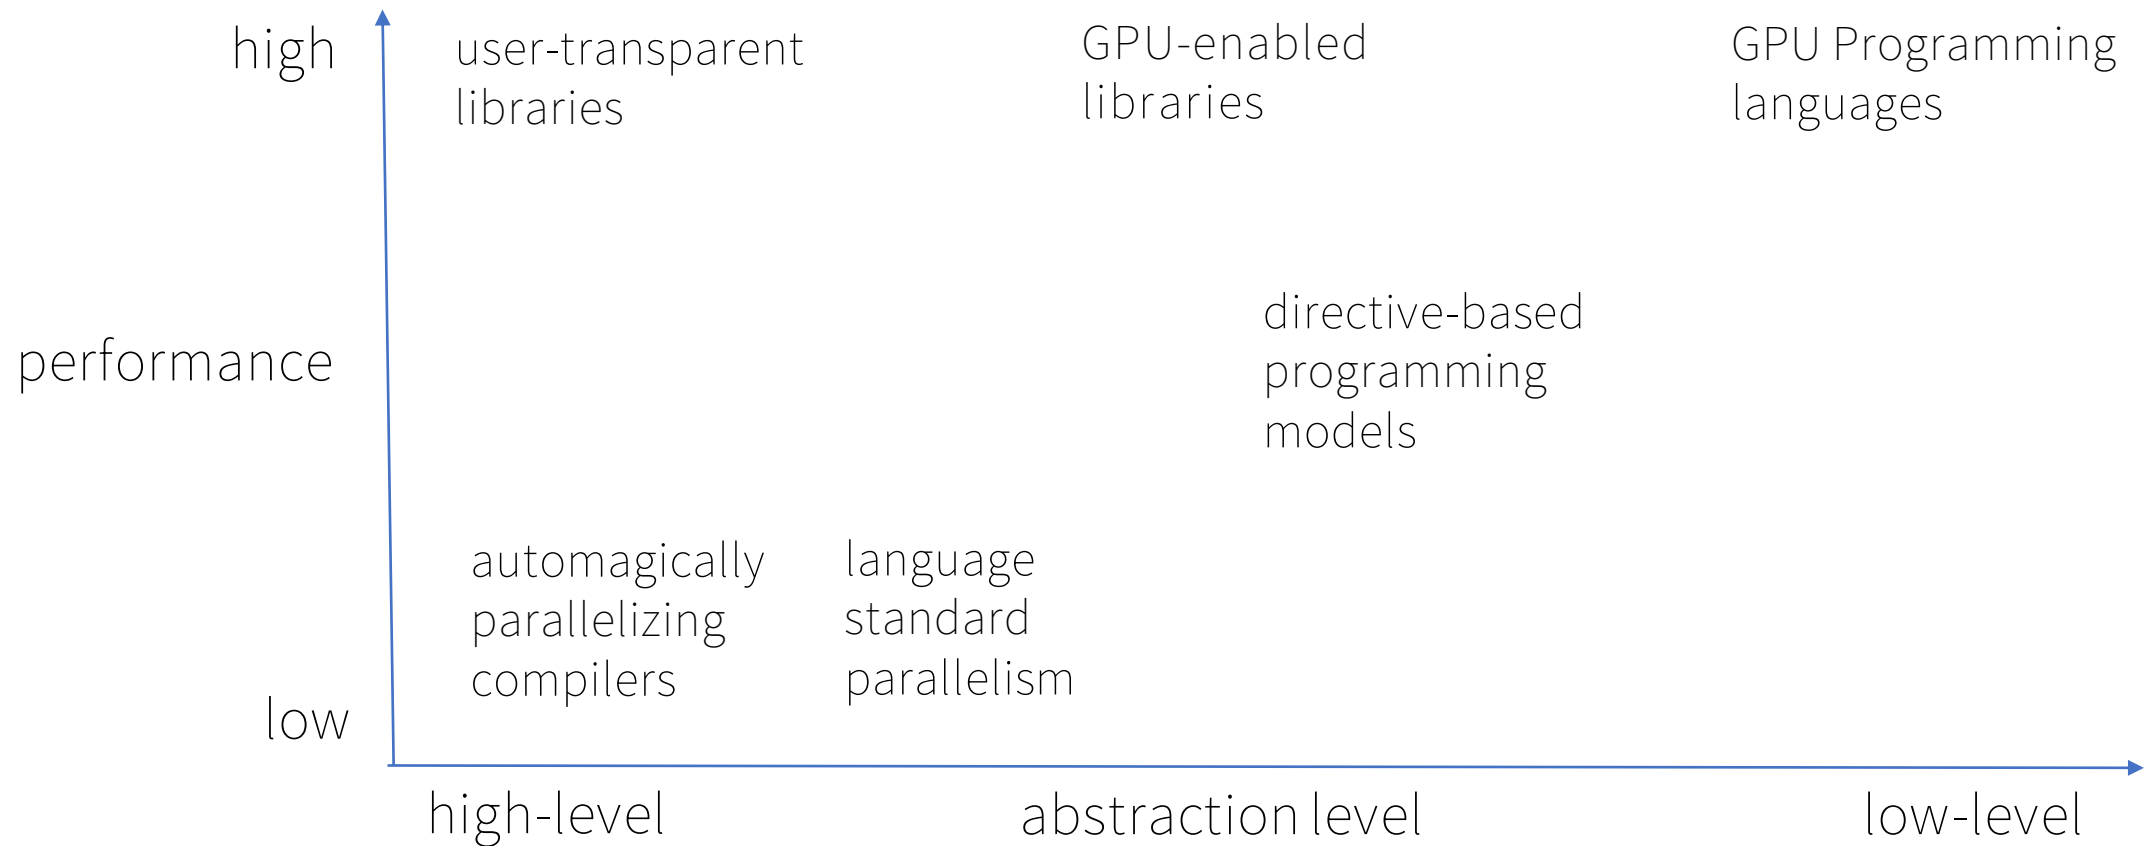
- Have to deal with separate memory space, limited bandwidth between host and device memory





GPU-enabled libraries





- Generally, the user is responsible for managing GPU memory
- Often use specialized objects that represents data in GPU memory
- Easy access to highly-optimized and tuned GPU routines
- Either focused on specific functionality or offering a ‘GPU Array’-like datatype

- Fast Fourier Transforms: cuFFT, clFFT, hipFFT, rocFFT, vkFFT
 - BLAS (linear algebra): cuBLAS, clBlas, rocBlas, clBlast (auto-tuned), SYCL-BLAS
 - Random number generation: cuRAND, rocRAND
 - Sparse matrix operations: cuSparse, hipSparse
 - Deep neural networks: cuDNN, OpenCV DNN, SYCL-DNN
-
- Practically all of these can be used directly from C++, many have Python bindings, bindings for other languages are not that commonly available or only supported by relatively small open-source projects

- Matlab: gpuArray
 - Provides access to many operations using cuBLAS, cuFFT underneath
 - Also JIT-compiles groups of pointwise array operations into CUDA kernels
- Python
 - CuPy: A NumPy-compatible array library accelerated by CUDA
 - Includes functionality for compiling 'raw' CUDA kernels
 - Includes bindings to cuBLAS, cuFFT, cuDNN, ...
 - PyTorch: Open source machine learning framework
 - Includes Tensor data type with CUDA backend
 - Can be used to interface cuBLAS, cuRAND, cuFFT, cuDNN, cuSPARSE, ...
 - Numba: Open source JIT compiler for Python/NumPy code
 - Compiles to CUDA or ROCm
 - Includes Python bindings to CUDA
 - cuDF: A Pandas-like GPU DataFrame library
 - Supports operations for loading, joining, aggregating, filtering, and otherwise manipulating data
 - Integrates with Dask for distributed and out-of-core computations
- ArrayFire, can be used from C, Rust, or Python
 - Includes functions for many image processing, linear algebra, and machine learning operations

CUDA SDK Samples include many simple example codes to illustrate how to use cuBLAS, cuFFT, and so on: <https://developer.nvidia.com/cuda-code-samples>

Examples on how to use libraries from AMD's ROCm platform are included in the documentation: <https://rocmdocs.amd.com>

Intel's [OneAPI toolkit](#) includes many numerical libraries and code samples

The datatype-oriented libraries often include their own memory managers, which can be great but sometimes complicates interoperability.

If you are not using C++ or Python, it is generally possible to write a small C++ code that calls the library function and can be called from another language, e.g. Fortran.

Introduction to CUDA Programming



Before we start:

- We will explain the CUDA Programming model
- We'll try to avoid talking about the hardware for now
- For the moment, please make no assumptions about the backend or how the program is executed by the hardware
- I will be using the term 'thread' a lot, this stands for '*thread of execution*' and should be seen as a parallel programming concept. Do not compare them to CPU threads.

The CUDA programming model separates a program into a host (CPU) and a device (GPU) part.

The host part:

- Allocates memory and transfers data between host and device memory, and starts GPU functions

The device part:

- Consists of functions that execute on the GPU, which are called *kernels*
- Kernels are executed by huge amounts of threads at the same time
- The data-parallel workload is divided among these threads
- The CUDA programming model allows you to code for each thread individually

- Parallelizing a computation sometimes requires to rethink algorithms, for example:

```
//some sort of stencil, reads 'a' once, 3 writes to a_new
for (int i=1; i<N-1; i++) {
    double my_a = a[i];
    a_new[i-1] += 0.25*my_a;
    a_new[i]   += 0.5*my_a;
    a_new[i+1] += 0.25*my_a;
}
```




```
//almost the same, but with 3 reads for every 1 write
for (int i=1; i<N-1; i++) {
    a_new[i] = 0.25*a[i-1] + 0.5*a[i] + 0.25*a[i+1];
}
```

The latter is much easier to parallelize because it avoids concurrent writes to the same memory locations

- The programming language for kernels is CUDA. It's mostly C/C++, but with some additions and limitations
- Additions:
 - Function qualifiers **__global__**, **__device__**, and **__host__** can be used to declare a function as being a kernel, a device function, or a host function
 - Kernel and device functions have built-in variables, like **threadIdx.xyz** or **blockIdx.xyz**
 - Memory qualifiers **__constant__** and **__shared__** can be used to declare a variable to reside in special memory spaces
 - Many intrinsic functions, e.g. **__sincosf()**, exist to use special function units in the hardware
- Limitations:
 - You cannot use any existing C functions, only functions with the **__device__** qualifier can be called from kernels.
 - A lot of standard C library functionality is not present, for example there is no **malloc()**, and for the first couple of years of CUDA there wasn't even a **printf()** function

- Kernels are executed in parallel by possibly millions of threads, so it makes sense to try to organize them in some manner



- In the CUDA programming model a thread is the most fine-grained entity that performs computations
- Threads within a kernel all execute the same program
- Threads direct themselves to different parts of memory using their built-in variables **threadIdx.xyz** (thread index *within* the thread block)

- Example:

```
for (i=0; i<N; i++) {  
    c[i] = a[i] + b[i];  
}
```

Create a single thread block of N threads:

```
i = threadIdx.x;  
c[i] = a[i] + b[i];
```

- Effectively the loop is ‘unrolled’ and spread across N threads

- Threads are grouped in thread blocks, allowing you to work on problems larger than the maximum thread block size
- Thread blocks are also numbered, using the built-in variable **blockIdx.xy** containing the index of each block within the grid.
- Total number of threads created is always a multiple of the thread block size, possibly not exactly equal to the problem size
- Other built-in variables are used to describe the thread block dimensions **blockDim.xyz** and grid dimensions **gridDim.xy**

- The host program sets the number of threads and thread blocks when it launches the kernel

```
// create variables to hold grid and thread block dimensions
```

```
dim3 threads(x, y, z);
```

```
dim3 grid(x, y, z);
```

```
// launch the kernel
```

```
vector_add<<<grid, threads>>>(c, a, b);
```

```
// wait for the kernel to complete
```

```
cudaDeviceSynchronize();
```

- The host program sets the number of threads and thread blocks when it launches the kernel

```
# create variables to hold grid and thread block dimensions
```

```
threads = (x, y, z)
```

```
grid = (x, y, z)
```

```
# launch the kernel
```

```
vector_add([c, a, b], block=threads, grid=grid)
```

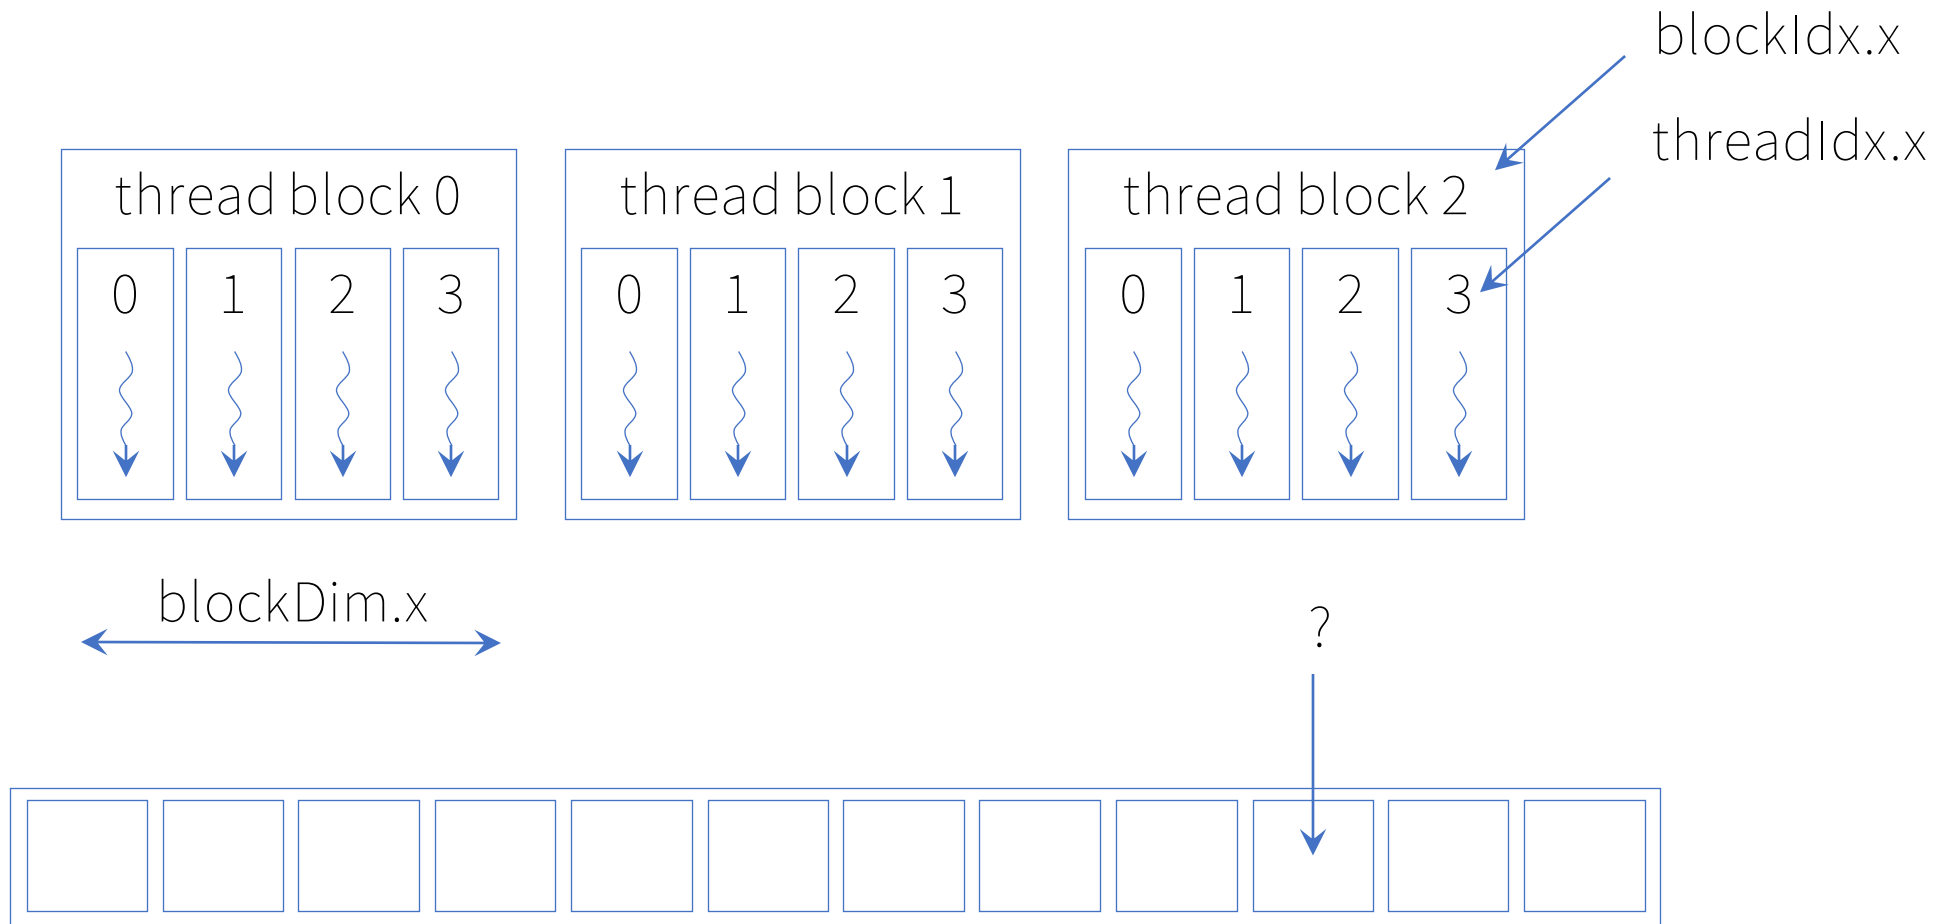
```
# wait for the kernel to complete
```

```
context.synchronize()
```

First Hands-on



- Connect to JupyterHub
 - <https://jupyter.lisa.surfsara.nl/jhlsrf021>
- Select the notebook for the first exercise
 - vector_add.ipynb
- Make sure you understand everything in the code, and complete the exercise!
- Hints:
 - Look at how the kernel is launched in the host program
 - <https://documen.tician.de/pycuda/driver.html#pycuda.driver.Function>
 - **threadIdx.x** is the thread index within the thread block
 - **blockIdx.x** is the block index within the grid
 - **blockDim.x** is the dimension of the thread block




Cuda Memories (part 1)

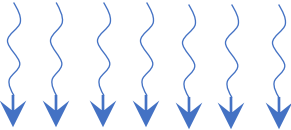


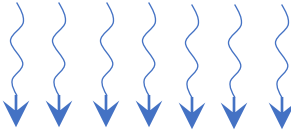
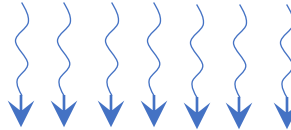
Registers

Shared memory

Global memory
Constant memory

Thread


Thread
Block


Grid	
(0, 0)	(1, 0)
	

- Example:

```
__global__ void matmul_kernel(float *C, float *A, float *B) {  
    int tx = threadIdx.x;        //local variable in registers  
    float local_sum[4];          //small compile-time sized array in registers
```

- Registers
 - Thread-local scalars or small constant size arrays are stored as registers
 - Implicit in the programming model
 - Behavior is very similar to normal local variables
 - Not persistent, after the kernel has finished, values in registers are lost

- Example:

```
__global__ void matmul_kernel( float *C,  //C points to global memory
                               float *A,  //A points to global memory
                               float *B)  //B points to global memory
```

- Global memory
 - Allocated by the host program using **cudaMalloc()**
 - Initialized by the host program using **cudaMemcpy()** or previous kernels
 - Persistent, the values in global memory remain across kernel invocations
 - Not coherent, writes by other threads will not be visible until kernel has finished

```
__constant__ float filter[filter_width * filter_height]; //initialized by a host

__global__ void convolution_kernel(float *output, float *input) {
    ...
    for (j = 0; j < filter_height; j++) {
        for (i = 0; i < filter_width; i++) {
            sum += input[y + j][x + i] *
                filter[j * filter_width + i]; //j and i do not depend on x and y
        }
    }
}
```

- Constant memory
 - Statically defined by the host program using **__constant__** qualifier
 - Defined as a global variable, visible only within the same translation unit
 - Initialized by the host program using **cudaMemcpyToSymbol()**
 - Read-only to the GPU, cannot be accessed directly by the host
 - Values are cached in a special cache optimized for broadcast access by multiple threads simultaneously, access should not depend on **threadIdx**

Second hands-on



- Connect to JupyterHub
 - <https://jupyter.lisa.surfsara.nl/jhlsrf021>
- Select the notebook for the second exercise
 - pnpoly.ipynb
- Make sure you understand everything in the code, and complete the exercise!
- Hints:
 - Use constant memory instead of global memory for the list of vertices
 - Python users can use `memcpy_htod()`, but need to find the symbol to copy to
 - See [PyCuda documentation on `get_global`](#)


Cuda Memories (part 2)

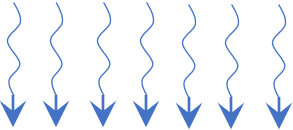


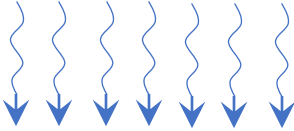
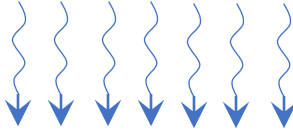
Registers

Shared memory

Global memory
Constant memory

Thread


Thread
Block


Grid	
(0, 0)	(1, 0)
	

```
__global__ void histogram(int *output, int *values, int n) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    __shared__ int sh_output[NUM_BINS];                //declare shared memory array

    if (threadIdx.x < NUM_BINS)
        sh_output = 0;                                //initialize shared memory
    __syncthreads();                                   //wait for all threads

    if(i < n) {
        int bin = values[i];
        atomicAdd(&sh_output[bin], 1);                 //increment bin in shared memory
    }
    __syncthreads();                                   //wait for all threads
    ...
}
```

- Shared memory
 - Variables have to be declared using **__shared__** qualifier, size known at compile time
 - In the scope of a thread block, all threads in a thread block see the same piece of memory
 - Not initialized, threads have to fill shared memory with meaningful values
 - Not persistent, after the kernel has finished, values in shared memory are lost
 - Not coherent, **__syncthreads()** is required to make writes visible to other threads within the thread block

```
__global__ void transpose(int h, int w, float* output, float* input) {
    int i = threadIdx.y + blockIdx.y * block_size_y;
    int j = threadIdx.x + blockIdx.x * block_size_x;

    __shared__ float sh_mem[block_size_y][block_size_x];    //declare shared memory array


    if (j < w && i < h) {
        sh_mem[threadIdx.y][threadIdx.x] = input[i*w+j];    //fill shared with values from global
    }
    __syncthreads();    //wait for all threads in the block

    i = threadIdx.x + blockIdx.y * block_size_y;
    j = threadIdx.y + blockIdx.x * block_size_x;
    if (j < w && i < h) {
        output[j*h+i] = sh_mem[threadIdx.x][threadIdx.y];    //store to global using shared memory
    }
}
```

Third hands-on

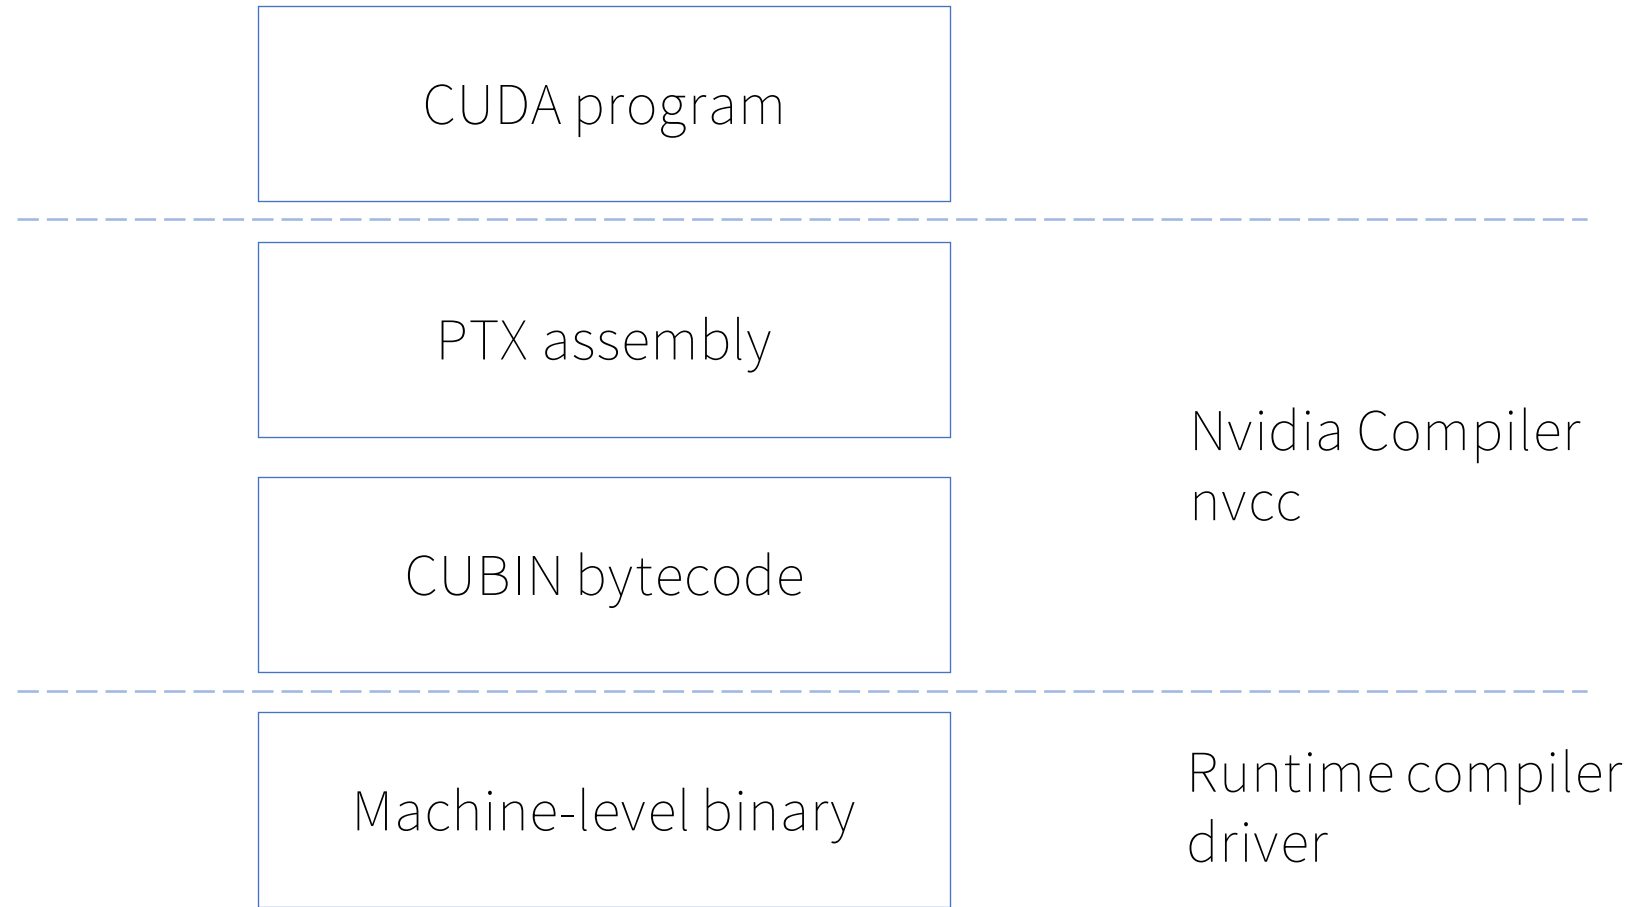


- Connect to JupyterHub
 - <https://jupyter.lisa.surfsara.nl/jhlsrf021>
- Select the notebook for the second exercise
 - reduction.ipynb
- Implement the kernel such that shared memory is used to sum the per-thread partial sums into a single per-thread block partial sum
- Make sure you understand everything in the code, and complete the exercise!
- Hints:
 - The number of thread blocks does not depend on n. All threads from all blocks first iterate (collectively) over the problem size (n) to obtain a per-thread partial sum
 - Within the thread block the per-thread partial sums are to be combined to a per-thread block partial sum
 - Each thread block stores its partial sum to **out_array[blockIdx.x]**
 - The kernel is called twice, the second kernel is executed with only one thread block to combine all per-block partial sums to a single sum



CUDA Program execution





CUDA	OpenCL	OpenACC	OpenMP 4
Grid	NDRange	compute region	parallel region
Thread block	Work group	Gang	Team
Warp	CL_KERNEL_PREFERRED_WORK_GROUP_SIZE_MULTIPLE	Worker	SIMD Chunk
Thread	Work item	Vector	Thread or SIMD

- Note that the mapping is implementation dependent for the open standards and may differ across computing platforms
- Not too sure about the OpenMP 4 naming scheme, please correct me if wrong

- Remember: all threads in a CUDA kernel execute the exact same program
- Threads are actually executed in groups of (32) threads called *warps*
- Threads within a warp all execute one common instruction simultaneously
- The context of each thread is stored separately, as such the GPU stores the context of all currently active threads
- The GPU can switch between warps even after executing only 1 instruction, effectively hiding the long latency of instructions such as memory loads

- All threads in a warp execute the exact same *instruction* at the same cycle

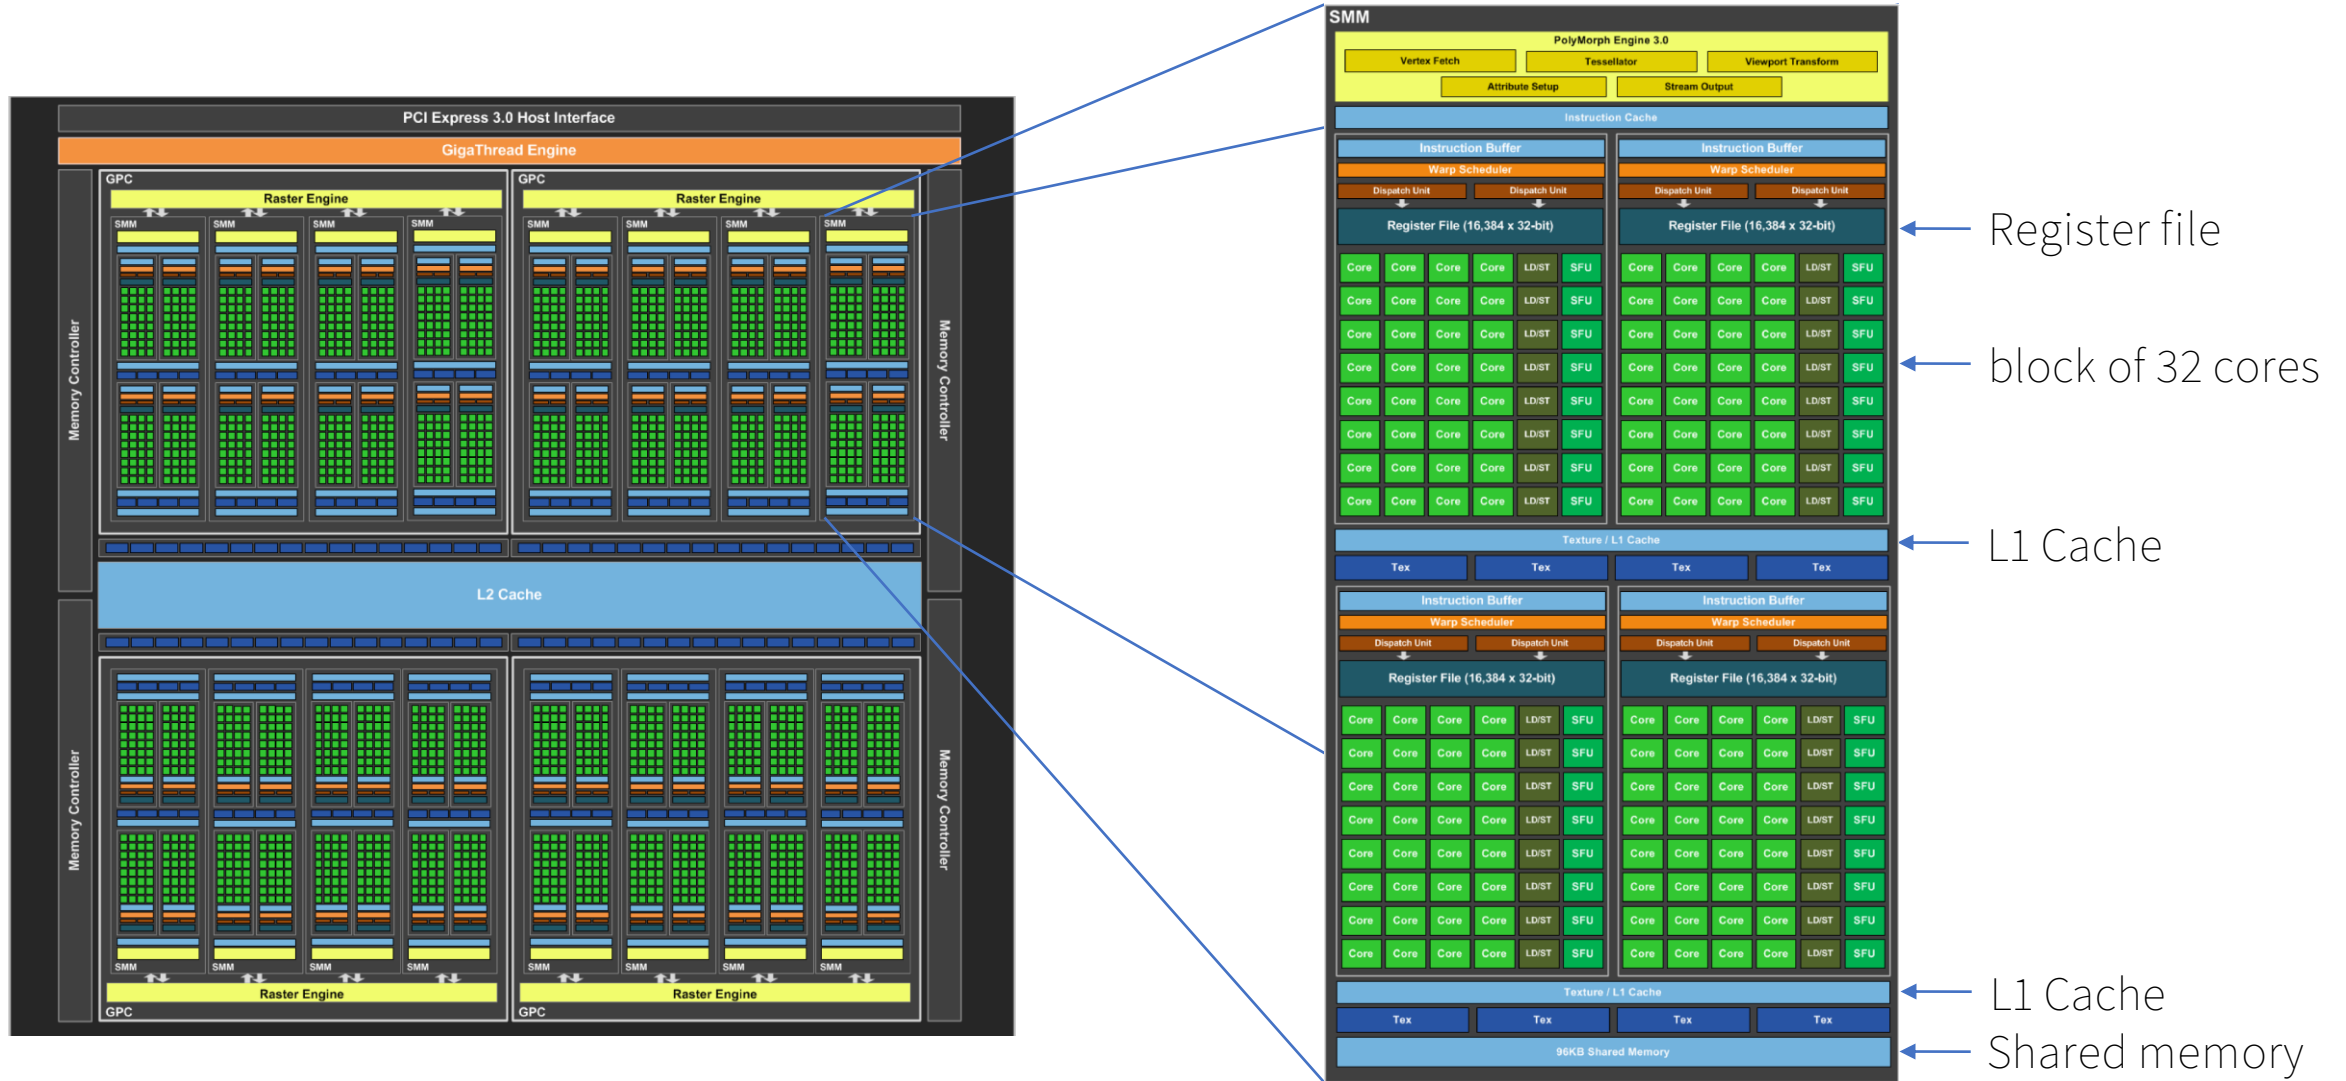
```
mad.f32    %f1, %f2, %f3, %f1;          // c += a*b;
```

- The same instruction, but on different data
- What about control flow instructions? (if, else, for, while)
 - All threads in the warp execute all live paths, with some threads predicated

```
if (a > 0.0f)
```
 - This is less efficient, but not always bad.
 - Avoid data-dependent conditional branching if possible
- Thread index-dependent branching is usually harmless, in particular when you respect the warp size

```
if (threadIdx.x < 32)
```
- The Volta architecture replaces predication with a per-thread program counter and call stack. The same performance recommendations apply, however.

Maxwell Architecture

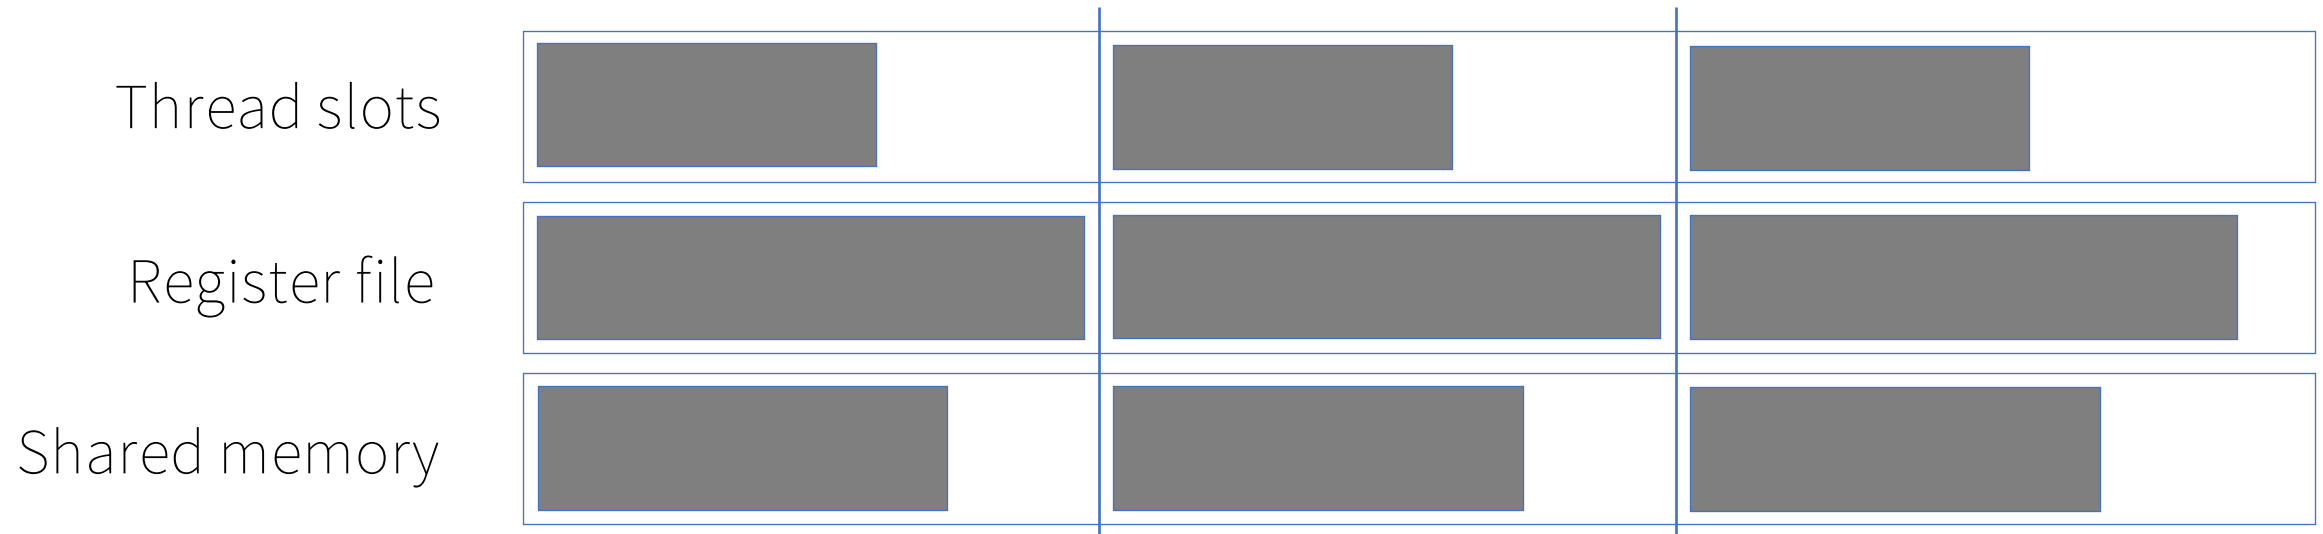


Turing architecture


- Features specialized Tensor and RT cores
- Tensor cores can operate on 4/8/16 bit integers and 16 bit half-precision floating points
- RT cores used for Ray-Tracing in graphics workloads



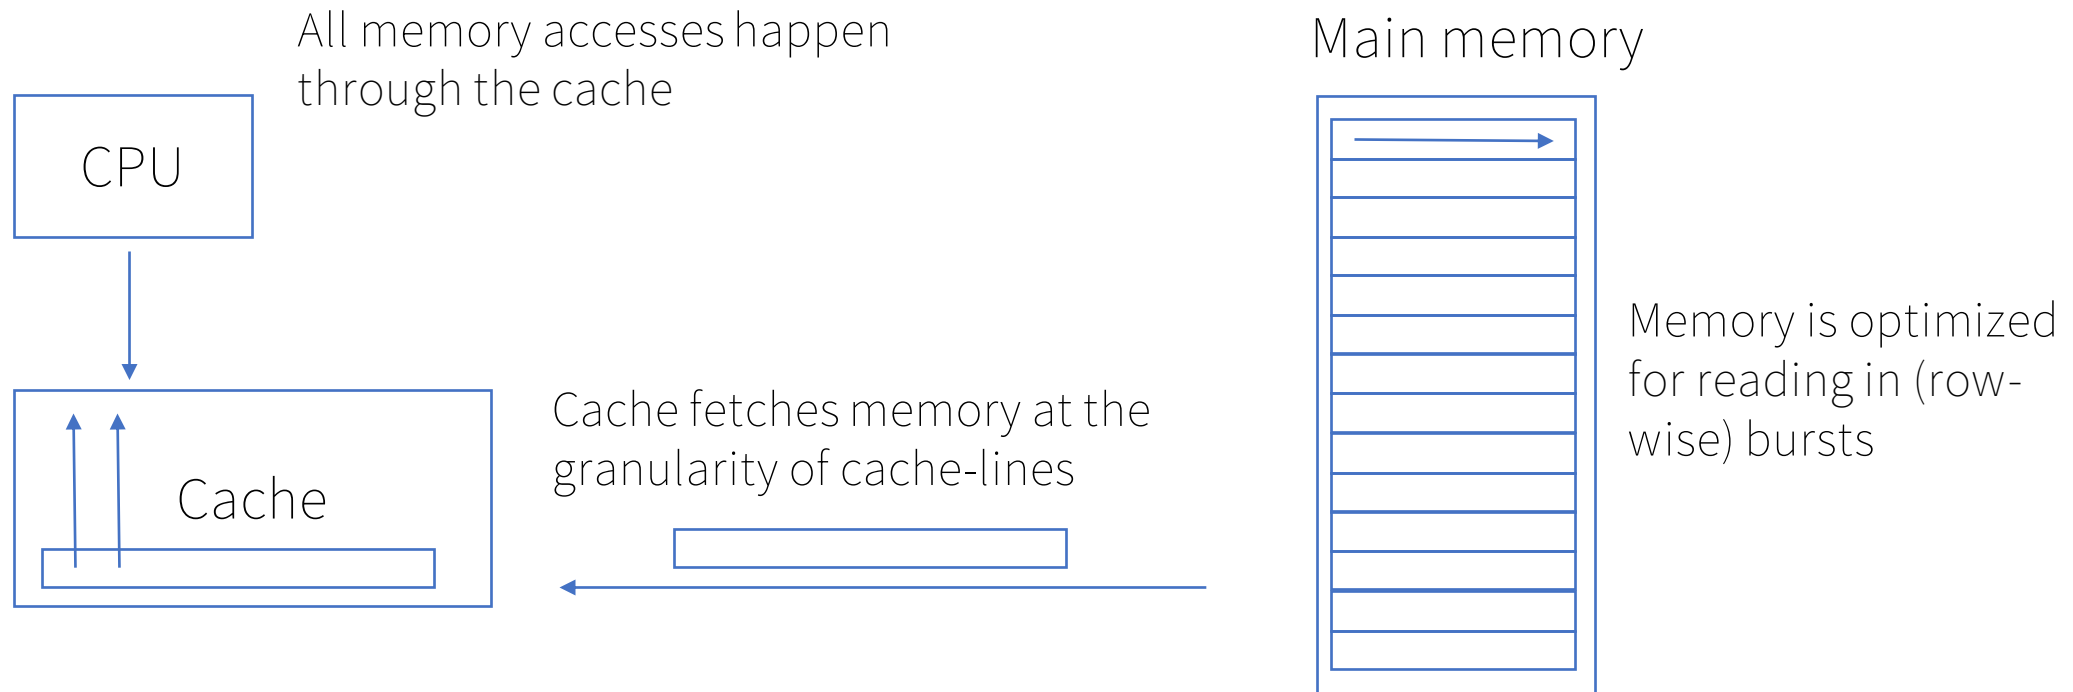
- The GPU consists of several (1 to 80) *streaming multiprocessors* (SMs)
- The SMs are fully independent
- Each SM contains several resources: Register file, Shared memory, Thread Slots, and Thread Block slots
- SM resources are dynamically partitioned among the thread blocks that execute concurrently on the SM, resulting in a certain *occupancy*



- Global memory is cached at L2, and for some GPUs also in L1
- When a thread reads a value from global memory, think about:
 - The total number of values that are accessed by the warp that the thread belongs to
 - The cache line length and the number of cache lines that those values will belong to
 - Alignment of the data accesses to that of the cache lines

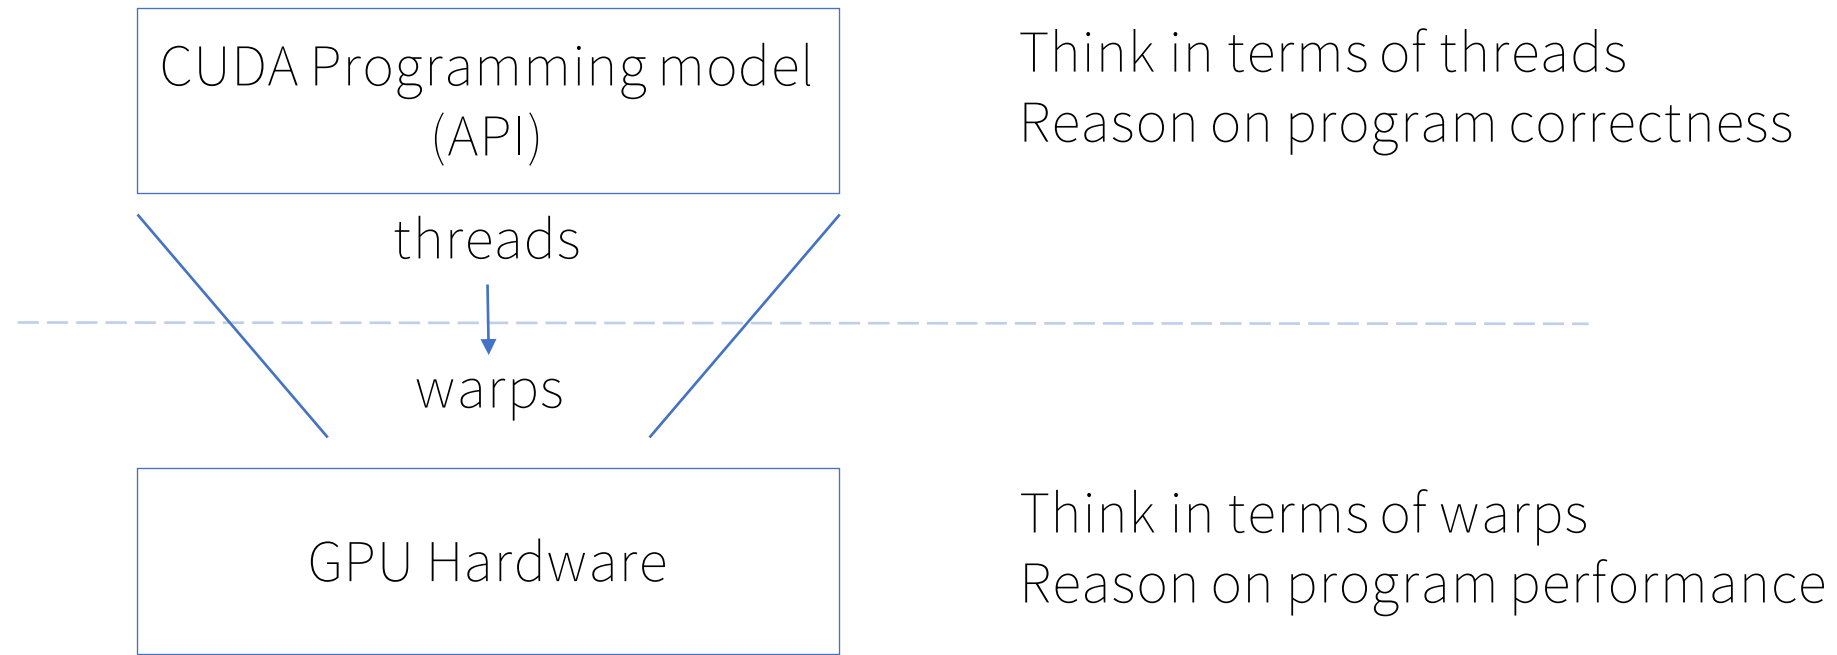


The memory hierarchy is optimized for certain access patterns



Subsequently accessing values that are adjacent on the same cache line is much faster than when each access requires a new cache line to be fetched

- Moving data around is more expensive than computing on it
- Start with a simple algorithm and keep it for readability and correctness checks
- Optimize only when needed
- Focus on the bottlenecks first
- Auto-tune (automatically explore the parameter space)
 - Different loop orderings
 - Different tile sizes, on multiple levels L3, L2, and L1
 - Different number of threads, thread blocks, vector lengths, etc
 - e.g. using Kernel Tuner (https://github.com/KernelTuner/kernel_tuner)



Closing remarks




- Read these slides again
 - And follow the links to the documentation!
- Play with the notebooks
 - If you do not have access to GPUs, try [Google Colab](#), they have GPU support
- Check the C code for the exercise
- If you have questions, contact us
 - b.vanwerkhoven@esciencecenter.nl
 - a.sclocco@esciencecenter.nl

- *Optimization Techniques for GPU Programming*
P. Hijma, S. Heldens, A. Sclocco, B. van Werkhoven, and H.E. Bal
ACM Computing surveys 2022
<https://dl.acm.org/doi/abs/10.1145/3570638>
- *Lessons Learned in a Decade of Research Software Engineering GPU Applications*
B. van Werkhoven, W.J. Palenstijn, A. Sclocco
International Conference on Computational Science (ICCS 2020)
https://doi.org/10.1007/978-3-030-50436-6_29
- *Kernel Tuner: A search-optimizing GPU code auto-tuner*
B. van Werkhoven
Future Generation Computer Systems 2019
<https://doi.org/10.1016/j.future.2018.08.004>

ESiWACE - for future exascale weather and climate simulations

Through the ESiWACE2 project we provide services on porting and performance optimization for weather and climate models

www.esiwace.eu




Collaborate with us!

We collaborate with researchers from universities and research institutes across the Netherlands on projects in every major discipline.

We also participate in many EU-funded Horizon 2020 projects.

www.esciencecenter.nl



Funding

The CORTEX project has received funding from the Dutch Research Council (NWO) in the framework of the NWA-ORC Call (file number NWA.1160.18.316).

ESiWACE2 has received funding from the European Union's Horizon 2020 research and innovation programme under grant agreement No 823988.



The ConFu project is funded by the Netherlands eScience Center (file number 00020223-A).